

# Coding Assignment:

## Fully Connected Neural Net (FCNN)

### Deadline: February 08, 2026

**Objective:** Build a Deep Learning engine from the ground up, then scale to industrial frameworks.

---

### Part 1: The "No-Framework" Challenge (NumPy Only)

Dataset: [UCI Adult Census Income](#)

*Goal: Understand the "Black Box" by building it.*

#### Question 1.1: Building the Engine

Using only **NumPy** and standard Python libraries (NO PyTorch/TensorFlow/Keras), implement a 3-layer FCNN.

- **The Task:** Manually code the following:
  - **Layer Initialization:** He-Initialization for weights and zeros for biases.
  - **The Forward Pass:**  $Z = W * X + b$ , followed by a ReLU activation.
  - **The Optimizer:** Implement the basic SGD update rule:

$$W_{\text{new}} = W_{\text{old}} - \text{learning\_rate} * \partial L / \partial W.$$

- **The Backpropagation:** Manually derive and code the chain rule for the gradients of weights and biases.
- **Deliverable:** A training loop that outputs the loss every 100 iterations. Show that your "From-Scratch" model can achieve at least 75% accuracy on the Census test set.

#### Question 1.2: The Importance of "Pre-Processing"

- **The Task:** Train your NumPy model twice: [once with raw data](#) and [once with Min-Max Scaled data](#).
- **Deliverable:** Compare the two results. Document what happens to the gradients when features have vastly different scales (e.g., "Age" vs. "Capital Gain").

---

## Part 2: Vision & Feature Interpretation (Frameworks Allowed)

**Dataset:** MNIST Handwritten Digits ([download “MINIST.zip” from Moodle](#))

*Goal: Bridge the gap between pixels and patterns using PyTorch/TensorFlow.*

### Question 2.1: Weight Visualization

- **The Task:** Build an FCNN using a library (PyTorch/TesorFlow). Once trained, extract the weights of the first hidden layer.
- **Deliverable:** Reshape the weights of 10 different neurons back into “28 x 28” grids and visualize them as heatmaps.
- **Observation:** Describe the "shapes" the network is looking for. Are they dots, lines, or just noise?

### Question 2.2: The "Flattening" Experiment

- **The Task:** Randomly shuffle the pixels of every image in the MNIST dataset (use the same shuffle pattern for all images). Train your FCNN on this "scrambled" data.
  - **Deliverable:** Compare the accuracy of the model on "Normal MNIST" vs. "Scrambled MNIST."
  - **Observation:** Why does the FCNN perform almost identically on both, while a human would find the scrambled version impossible? (This highlights the lack of spatial awareness in FCNNs).
- 

## Part 3: Stress Testing & Robustness

**Dataset:** Tiny ImageNet 10 - a subset of tiny-imagenet-200 available at

<https://www.kaggle.com/datasets/akash2sharma/tiny-imagenet> ([Download “tiny-imagenet.zip” from Moodle](#))

### Question 3.1: Vanishing Gradients & Modern Fixes

- **The Task:** Create a "Very Deep" FCNN (8+ layers).
  - **Experiment A:** Use **Sigmoid** activations throughout.
  - **Experiment B:** Use **ReLU** activations + **Batch Normalization**.
- **Deliverable:** Plot the "Gradient Norm" (magnitude of the weights' updates) for the first layer in both experiments. Explain which version trained faster and why.

### Question 3.2: The Ablation Study

Perform a "Switch-Off" test. Report how the final accuracy changes when you:

1. Remove **Dropout**.
2. Change the **Learning Rate** by a factor of 10 (too high vs. too low).

3. Switch from **Adam** to **Vanilla SGD**.
  - **Deliverable:** A summary table showing which factor had the biggest impact on performance
- 

## Submission Requirements

1. **Code:** Submit a well-documented Jupyter Notebook (**.ipynb**).
2. **Report** (**.pdf**) detailing:
  - Problem definition & methodology.
  - Hyperparameter tuning & architecture choices.
  - Results, visualizations, and interpretations.
  - Key findings from latent space analysis.

### Guidelines for the Submission:

1. Upload the Assignment Report in PDF format with the following name:  
**<Group\_number>\_Assignment1\_Report.pdf**
2. Upload the code files in a single zip file with the following name:  
**<Group\_number>\_Assignment1\_Code.zip**

\*\*\*\*\*